

PAR – In-Term Exam – Course 2022/23-Q1

November 3rd, 2022

Problem 1 (3.0 points) Given the following code:

```
#define N 256
#define BS 64
int m[N][N];

for (int ii=0; ii<N/BS; ii++) {
    for (int jj=0; jj<N/BS; jj++) {
        tareador_start_task("task");

        // In general, a task processes a block of BSxBS elements
        // However, if the task is labeled with (ii,jj=0),
        // this task processes BSx(BS-1) elements
        int i_start = ii*BS; int i_end = i_start+BS;
        int j_start = jj*BS; int j_end = j_start+BS;
        for (int i=i_start; i<i_end; i++)
            for (int j=max(j_start,1); j<j_end; j++)
                m[i][j] = compute (m[i][j], m[i][j-1]); // tc t.u. (time units)

        tareador_end_task("task");
    } }
```

Assume the innermost loop body takes t_c t.u., all variables but matrix m are in registers, function `compute` only reads the values received as arguments and does not modify other positions in the memory, BS perfectly divides N , being BS and N defined in the code above. **We ask you:**

1. (1.0 points) Draw the task dependence graph (TDG), indicating the cost of each task as a function of BS and t_c . Label each task with the values of ii and jj .
2. (1.0 points) Compute T_1 , T_∞ , P_{min} as a function of BS , N and t_c .

3. (1.0 points) Assuming the assignment of tasks to 4 processors of the table below, calculate T_4 and speedup S_4 .

Processor	Tasks
P_0	$(0, 0), (1, 0), (2, 0), (3, 0)$
P_1	$(0, 1), (1, 1), (2, 1), (3, 1)$
P_2	$(0, 2), (1, 2), (2, 2), (3, 2)$
P_3	$(0, 3), (1, 3), (2, 3), (3, 3)$

Problem 2 (2.0 points) Given the same code and mapping of tasks to 4 processors as in the previous exercise, assume

- A distributed-memory architecture with $P = 4$ processors;
- Matrix m is initially distributed by columns (BS consecutive columns per processor).
- Data sharing model with $t_{comm} = t_s + W \times t_w$, being W the number of elements to transfer, and t_s and t_w the start-up time and transfer time of one element, respectively;
- The execution time for a single iteration of the innermost loop body takes t_c t.u.

We ask you: Draw the execution timeline of the execution of tasks and write the expression that determines the execution time T_P , clearly indicating the contribution of the computation time $T_{P_{comp}}$ and data sharing overhead $T_{P_{mov}}$, as a function of N , BS , P , t_c , t_s and t_w .

Problem 3 (5.0 points) Consider a multiprocessor system with a hybrid NUMA/UMA architecture which is composed by 3 identical NUMAnodes. Each NUMAnode has 20 Gbytes of main memory and 2 processors with its own private cache of 8 Mbytes. The memory cache lines are 32 bytes wide, and data coherence is guaranteed using *Write-Invalidate MSI protocol* within each NUMAnode and using a *Write-Invalidate MSU Directory-based* cache coherency protocol among NUMAnodes.

We ask you to answer the following questions:

1. (1.0 points) Compute the total number of bits that are necessary **in each cache memory** to maintain the coherence, indicating the function of those bits.
2. (1.0 points) Compute the total number of bits that are necessary **in each NUMAnode's directory** to maintain the coherence, indicating the function of those bits.

Given the following OpenMP code excerpt which is executed on processor 0 from the previous described multiprocessor system:

```
#define N (6*1024)
#define NUM_THREADS 6
int v[N], count[NUM_THREADS];
...
for (int k = 0; k < NUM_THREADS; k++)
    count[k]=0;

/**** POINT A ****/

#pragma omp parallel num_threads (NUM_THREADS)
{
    int id=omp_get_thread_num();
    int num_iter = N/NUM_THREADS;

    for (int k = id*num_iter; k < id*num_iter+num_iter; k++) {
        int value = v[k];
        if (is_prime(value)) /* returns true if "value" is a prime number */
            count[id]++;
    }
}
```

and assuming that: 1) the initial memory address of vector count is aligned to the start of a memory/cache line; 2) the size of an int data type is 4 bytes; and 3) processors 0 and 1 belong to NUMAnode0, processors 2 and 3 belong to NUMAnode1 and processors 4 and 5 belong to NUMAnode2; and 4) the Operating System applies the "first touch" policy for data allocation in memory. **We ask you to:**

3. (1.25 points) Complete the table in the provided answer sheet with the required information: affected cache line (numbered from the first position where vector count is allocated), cache line states (I/S/M) in processors 0 to 5, directory entry state (U/S/M) and presence bits (0/1, where the lowest ordered bit, the rightmost one, corresponds to NUMAnode0), to keep cache coherence, **when the execution of the previous code reaches POINT A.**

4. (1.25 points) Complete the table in the provided answer sheet with the required information: affected cache line, access in cache (hit/miss), CPU command for processor k ($PrRd_k/PrWr_k$), Bus transaction(s) from Snoopy in processor k ($BusRd_k / BusRdX_k / BusUpgr_k / Flush_k / Nothing$), cache line states in processors 0 to 5, directory entry state and presence bits, to keep cache coherence, **AFTER the execution of each** memory access in the table. Assume initially memory and cache states from previous question.

5. (0.5 points) Have you observed any potential efficiency problem in the previous code? Justify briefly your answer.

Student name:

Answer for question 3.3

	Affected line	Home NUMAnode	Cache line state						Directory entry	
			0	1	2	3	4	5	State	Presence bits
count[0]										
count[1]										
count[2]										
count[3]										
count[4]										
count[5]										

Answer for question 3.4

[illegible]