

# PAR – In-Term Exam – Course 2023/24-Q1

November 2<sup>nd</sup>, 2023

**Problem 1** (5.0 points) Given the following code:

```
#define BS 8
#define N 24
double A[N][N];
int ii, jj, i, j;
...

for (ii=0; ii<N; ii+=BS)
    for (jj=0; jj<N; jj+=BS) {
        tareador_start_task("Comp_ii_jj");
        for (i=max(1,ii); i<min(ii+BS,N); i++)
            for (j=max(1,jj); j<min(jj+BS,N-1); j++)
                A[i][j] = A[i][j-1] + A[i-1][j] + A[i][j+1]; // cost of this line: tc t.u.
        tareador_end_task("Comp_ii_jj");
    }
...
```

1. (1 point) Draw the Task Dependence Graph (TDG) based on the above Tareador task definitions and for BS=8 and N=24. Each task should be clearly labeled with the values of  $ii, jj$  and its cost in time units.
2. (2.0 points) Compute the values for  $T_1$ ,  $T_\infty$  and  $P_{min}$ . Draw the temporal diagram for the execution of the TDG in the previous question on  $P_{min}$  processors. As indicated, consider the cost of the innermost loop body to be  $t_c$  time units.

3. (2.0 points) Assume the same task definition, and consider a distributed memory architecture with  $P$  processors,  $BS = N/P$ , and  $N$  a very large value multiple of  $P$  (you can assume  $BS-1$  is approximately the same as  $BS$  to simplify the model). Let's assume that matrix  $A$  is initially distributed by columns ( $N/P$  consecutive columns per processor) and tasks are scheduled so that a task is executed in the processor that stores the data the task has to update.

**We ask you to:**

- (a) (1.0 points) Draw the time diagram for the execution of the tasks in  $P$  processors clearly identifying the computation and the data sharing time. Note: you can assume  $P = 3$  for this question a).
- (b) (1.0 points) Write the expression that determines the execution time,  $T_p$ , as a function of  $N$  and  $P$  clearly identifying the contribution of the computation time and the data sharing overheads, assuming the data sharing model explained in class in which the overhead to perform a remote memory access is  $t_s + t_w \times m$ , being  $t_s$  the start-up time,  $t_w$  the time to transfer one element and  $m$  the number of elements to be transferred; at a given time, a processor can only perform one remote access to another processor and serve one remote access from another processor.



**Problem 2** (5.0 points) Given the following sequential recursive algorithm:

```
#define N 1024
#define NSTATES 128
#define MINROWS 2

int histogram[NSTATES];

int base_processing (int data[N][N], int start, int nrows) {
    int outofrange=0;

    for (int i=start; i<start+nrows; i++) {
        for (int k=0; k<N; k++) {
            int value = compute (data[i][k]); /* perform computation on the parameter */
            if (value >= NSTATES)
                outofrange++;
            else if (value >= 0)
                histogram[value]++;
        }
    }
    return outofrange;
}

int rec_processing (int data[N][N], int start, int nrows) {
    int res1, res2=0;

    if (nrows < MINROWS)
        res1 = base_processing (data, start, nrows);
    else {
        res1 = rec_processing (data, start, nrows/2);
        res2 = rec_processing (data, start+nrows/2, nrows-nrows/2);
    }
    return res1 + res2;
}

int main() {
    int data[N][N];
    ...
    int res = rec_processing(data, 0, N);
    ...
}
```

**We ask you to** answer the following **independent** questions:

1. (2.5 points) Write an OpenMP parallel version of the `base_processing` function, following an *Iterative* Task Decomposition, making use of the OpenMP explicit tasks. Your implementation should minimize synchronization overheads and take into account **the potential imbalance** generated by the compute function.

Some considerations about the all the proposed solutions:

- All of them have `parallel` and `single`.
- All of them create explicit tasks.
- All of them avoid creating tasks with only one iteration of  $k$  (too much task creation overhead).
- All of them avoid creating tasks doing full loop  $k$  (too much imbalance due to compute).
- All of them avoid waiting for all tasks at each iteration of  $i$  (we are looking for parallelism). I.e. `nogroup` should be used in the taskloop, but then reduction is not allowed.
- All of them perform a reduction of `outofrange` variable (avoid data race condition and reduce data synchronization).
- All of them perform an `atomic` to update histogram (reduction could be possible also but it may be costly in memory).











2. (2.5 points) Write an OpenMP parallel version of the sequential recursive program, following a *Recursive Task Decomposition* using the *Tree* strategy. The implementation should take into account the overhead due to task creation, by limiting their creation once a certain level in the recursive tree is reached.

